

Distributed Filesystems

☰ Tags	CompSci Low-Level
⚙ Status	In progress
▼ Type	Topic

Goals

1. **Build a Mini Distributed File Store Simulator** in Rust within 10 consecutive days.
2. **Implement key features:** chunking, SHA-256 hashing, lossless compression (LZ4 or Huffman), configurable replication, XOR parity recovery, and a CLI.
3. **Exercise Rust fundamentals:** ownership, borrowing, lifetimes, Arc / Mutex , threads or async tasks, Send/Sync traits.
4. **Benchmark and document** throughput, storage overhead, and areas for optimisation.
5. **Produce artifacts:** public Git repo, release binary, README with usage examples, and a short retrospective.

Resources

Type	Title / Tool	Purpose
Paper	Google File System (2003); HDFS Architecture Guide	Storage design reference
Paper	Colossus/Spanner Papers (summaries)	Modern object-store inspiration
Book	<i>Designing Data-Intensive Applications</i> – ch. 2, 10	Distributed-file-system theory
Book	<i>The Rust Programming Language</i> – ch. 4, 15, 16	Ownership, lifetimes, concurrency

Crate	<code>clap</code> , <code>sha2</code> , <code>lz4_flex</code> , <code>reed-solomon-erasure</code>	CLI, hashing, compression, parity
Crate	<code>tokio</code> or <code>std::thread</code> + <code>crossbeam</code>	Async or thread pools
Tool	<code>cargo flamegraph</code> , <code>hyperfine</code>	Profiling and benchmarking
Blog	MinIO Erasure Coding Docs	Parity layout ideas

Project

Day-by-Day Outline (25–30 min slots scale-able to 3–4 h)

Day	Focus & Checkpoints
1	Read GFS & HDFS papers; decide chunk size, replication, parity; initialise Git/Cargo project
2	Scaffold CLI (<code>dfs-sim</code>) with commands: <code>put</code> , <code>get</code> , <code>node ls</code>
3	Implement file-to-chunk splitter and SHA-256 hashes; unit tests
4	Add in-memory <code>Node</code> and round-robin replication; basic catalog
5	Integrate compression crate; measure ratios on sample files
6	Implement XOR parity per block set; simulate node failure & recovery
7	Build download path: verify checksum, decompress, reassemble file
8	Add simple benchmarks and logging; profile hot paths
9	Polish code, error handling, parallel transfers (threads or Tokio)
10	Write README usage section, performance results, retrospective, tag v0.1.0

Stretch Goals (fit as time allows)

- Swap internal calls to gRPC streaming APIs.
- Prototype consistent-hash placement.
- Replace XOR with Reed-Solomon erasure coding.
- Compare your simulator's API to MinIO's S3 layer.

How it should look like?

```
dfs-sim/
├─ Cargo.toml
├─ README.md    – quick-start + design notes
├─ src/
│   ├─ main.rs    – CLI
│   ├─ cli.rs
│   ├─ catalog.rs
│   ├─ node.rs
│   ├─ chunker.rs – split + SHA-256
│   ├─ compressor.rs
│   ├─ parity.rs
│   └─ util.rs
```

Run a local cluster

```
# three storage nodes
dfs-sim node --id n1 --listen 127.0.0.1:4001
dfs-sim node --id n2 --listen 127.0.0.1:4002
dfs-sim node --id n3 --listen 127.0.0.1:4003

# coordinator
dfs-sim coord \
  --nodes 127.0.0.1:4001,127.0.0.1:4002,127.0.0.1:4003 \
  --replication 2 --parity xor
```

Basic CLI

```
dfs-sim put ./large.iso --name iso/ubuntu.iso
dfs-sim get iso/ubuntu.iso --out ./restore.iso
```

If a node is offline, download still succeeds via XOR rebuild.